

KEY DISTRIBUTION: PKI and SESSION-KEY EXCHANGE

Mihir Bellare

UCSD

1

But who exactly are “Alice” and “Bob”?

A question of **identity** that is central to key distribution and its security.

Mihir Bellare

UCSD

3

The public key setting

Bob's secret key is $sk[B]$ and its associated public key is $pk[B]$. The public key setting **assumes** Alice is in possession of $pk[B]$.

$$\begin{array}{ccc} \text{Alice}^{pk[B]} & & \text{Bob} \\ C \xleftarrow{\$} \mathcal{E}_{pk[B]}(M) & \xrightarrow{C} & M \leftarrow \mathcal{D}_{sk[B]}(C) \\ \mathcal{V}_{pk[B]}(M, \sigma) & \xleftarrow{M, \sigma} & \sigma \xleftarrow{\$} \mathcal{S}_{sk[B]}(M) \end{array}$$

Now Alice can encrypt a message M under $pk[B]$ to get a ciphertext C that Bob can decrypt using $sk[B]$.

Bob can sign a message M using $sk[B]$ to get signature σ that Alice can verify using $pk[B]$.

But how does Alice get $pk[B]$?

Mihir Bellare

UCSD

2

Names are bound to people via societal means

The question of identity

Alice

Bob

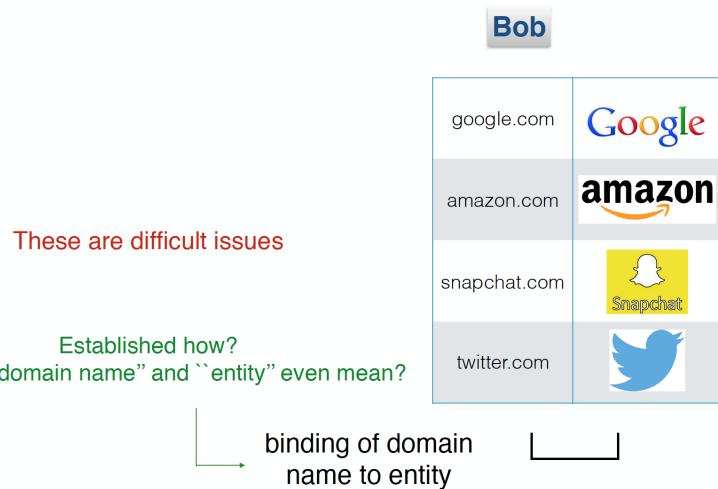


Mihir Bellare

UCSD

4

In TLS, Bob is a server



Mihir Bellare

UCSD

5

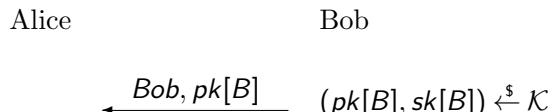
In TLS, Alice is a client

Alice's identity could be her ip address or domain name.

Alice need not be an individual; she could be a corporation, a server herself, ...

How can Alice get Bob's public key?

A simple idea: Bob generates his keys, and just sends $pk[B]$ to Alice with his identity “Bob” attached.



This enables:

$$\begin{array}{ccc} C \xleftarrow{\$} \mathcal{E}_{pk[B]}(M) & \xrightarrow{C} & M \leftarrow \mathcal{D}_{sk[B]}(C) \\ \mathcal{V}_{pk[B]}(M, \sigma) & \xleftarrow{M, \sigma} & \sigma \xleftarrow{\$} \mathcal{S}_{sk[B]}(M) \end{array}$$

Mihir Bellare

UCSD

7

Mihir Bellare

UCSD

8

PKI, CAs and certificates

Goal: Alice gets an **authentic** copy of Bob's public key, meaning if *pk* claims to come from Bob, she has proof to that effect.

Popular Solution: The PKI (Public Key Infrastructure).

Certificate authority: Trusted entity that provides the above proof.

Certificate: The proof

Note: There are other ways to reach the goal: Bob could post his public key on his Facebook; post it on his personal or corporate webpage; include it as an attachment in his emails; put it on a keyserver like openpgp SGS; hand it to Alice in person; ...

Mihir Bellare

UCSD

9

Let's Encrypt

Root Certificates

Our roots are kept safely offline. We issue end-entity certificates to subscribers from the intermediates in the next section.

- Active
 - [ISRG Root X1 \(self-signed\)](https://isrgrootx1.letsencrypt.org/)

We've set up websites to test certificates chaining to our roots.

- ISRG Root X1 Valid Certificate
 - <https://valid-isrgrootx1.letsencrypt.org/>
- ISRG Root X1 Revoked Certificate
 - <https://revoked-isrgrootx1.letsencrypt.org/>
- ISRG Root X1 Expired Certificate
 - <https://expired-isrgrootx1.letsencrypt.org/>

Intermediate Certificates

IdenTrust has cross-signed our intermediates. This allows our end certificates to be accepted by all major browsers while we propagate our own root.

Under normal circumstances, certificates issued by Let's Encrypt will come from "Let's Encrypt Authority X3". The other intermediate, "Let's Encrypt Authority X4", is reserved for disaster recovery and will only be used should we lose the ability to issue with "Let's Encrypt Authority X3". The X1 and X2 intermediates were our first generation of intermediates. We've replaced them with new intermediates that are more compatible with Windows XP.

- Active
 - [Let's Encrypt Authority X3 \(IdenTrust cross-signed\)](#)
 - [Let's Encrypt Authority X3 \(Signed by ISRG Root X1\)](#)

Mihir Bellare

UCSD

11

Let's Encrypt



Mihir Bellare

UCSD

10

Some certificate authorities

Rank	Issuer	Usage	Market share
1	IdenTrust	20.4%	39.7%
2	Comodo	17.9%	34.9%
3	DigiCert	6.3%	12.3%
4	GoDaddy	3.7%	7.2%
5	GlobalSign	1.8%	3.5%
6	Certum	0.4%	0.7%
7	Actalis	0.2%	0.3%
8	Entrust	0.2%	0.3%
9	Secom	0.1%	0.3%
10	Let's Encrypt	0.1%	0.2%
11	Trustwave	0.1%	0.1%
12	WISeKey Group	< 0.1%	0.1%
13	StartCom	< 0.1%	0.1%
14	Network Solutions	< 0.1%	0.1%

Mihir Bellare

UCSD

12

Certificate process

- Bob generates sk and pk by locally running key-generation
- Bob sends his identity and pk to CA
- CA does identity check to ensure pk is Bob's
- Bob proves knowledge of sk to CA
- CA issues certificate to Bob
- Bob sends certificate to Alice
- Alice verifies certificate and extracts Bob's public key pk

Mihir Bellare

UCSD

13

Key generation

Bob locally runs a prescribed key-generation algorithm to generate his keys: $(pk, sk) \xleftarrow{\$} \mathcal{K}$

Key generation example: RSA Key generation with openssl

Generating a private RSA key

1. Generate an RSA private key, of size 2048, and output it to a file named key.pem:

```
$ openssl genrsa -out key.pem 2048
Generating RSA private key, 2048 bit long modulus
.....+++
.....+++++
e is 65537 (0x10001)
```

2. Extract the public key from the key pair, which can be used in a certificate:

```
$ openssl rsa -in key.pem -outform PEM -pubout -out public.pem
writing RSA key
```

Mihir Bellare

UCSD

15

Key generation Example: EC Key generation with openssl

Generating a private EC key

1. Generate an EC private key, of size 256, and output it to a file named key.pem:

```
$ openssl ecparam -name prime256v1 -genkey -noout -out key.pem
```

2. Extract the public key from the key pair, which can be used in a certificate:

```
$ openssl ec -in key.pem -pubout -out public.pem
read EC key
writing EC key
```

After running these two commands you end up with two files: key.pem and public.pem. These files are referenced in various other guides on this page when dealing with key import.

Mihir Bellare

UCSD

16

Key generation Example: On-line key generation tools

The screenshot shows the "PGP Key Generator" interface. It includes fields for "Email Address" and "PGP Key Password / Passphrase". Below these are sections for "Generate PGP Keys" and "PGP Private Key" and "PGP Public Key". The "PGP Private Key" section contains a large block of PGP key data starting with "-----BEGIN PGP PRIVATE KEY BLOCK-----". The "PGP Public Key" section also contains a large block of PGP key data starting with "-----BEGIN PGP PUBLIC KEY BLOCK-----". Both sections include version information (BCPG C# v1.6.1.0) and detailed key parameters.

Mihir Bellare

UCSD

17

Checks

Bob sends his identity Bob (domain name, ip address, email address, ...) and his public key pk to the certificate authority (CA).

Upon receiving (Bob, pk) the CA performs some checks to ensure pk is really Bob's key.

Example: If Bob is a domain name, then the CA sends Bob a challenge and checks that he can put it on the webpage of the domain name.

Example: If Bob is an email address, then the CA sends an email to that address with a link for Bob to click to verify that he owns the address.

Example: If Bob is a passport or driver's license, the CA may be able to verify it physically, out of band.

Proof of knowledge of secret key: The CA might have Bob sign or decrypt something under pk to ensure that Bob knows the corresponding secret key sk . This ensures Bob has not copied someone else's key.

Mihir Bellare

UCSD

18

Certificate Issuance

Once CA is convinced that pk belongs to Bob , it forms a certificate

$$CERT_{Bob} = (CERTDATA, \sigma),$$

where σ is the CA's signature on $CERTDATA$, computed under the CA's secret key $sk[CA]$, and $CERTDATA$ contains:

- Bob's public key pk , and its type (RSA, EC, ...)
- Identity Bob of Bob
- Name of CA
- Expiry date of certificate
- ...

The certificate $CERT_{Bob}$ is returned to Bob.

Mihir Bellare

UCSD

19

Certificate usage

Bob can send $CERT_{Bob}$ to Alice who is assumed to have the CA's public key $pk[CA]$ and now will:

- Parse it as $(CERTDATA, \sigma) \leftarrow CERT_{Bob}$
- Check that $\nu_{pk[CA]}(CERTDATA, \sigma) = 1$
- Extract $(pk, Bob, expiry, \dots) \leftarrow CERTDATA$
- Check certificate has not expired
- ...

If all is well, Alice accepts the certificate and is ready to use the public key pk therein.

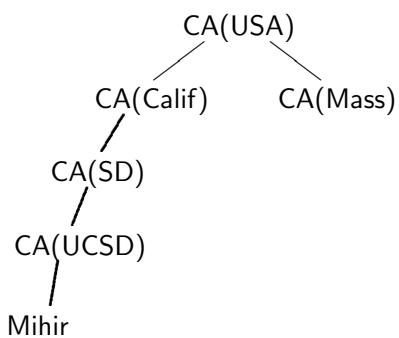
How does Bob get $pk[CA]$? CA public keys are embedded in software such as your browser.

Mihir Bellare

UCSD

20

Certificate hierarchies



CERT_{Mihir}

CERT[CA(USA) : CA(Calif)]
CERT[CA(Calif) : CA(SD)]
CERT[CA(SD) : CA(UCSD)]
CERT[CA(UCSD) : Mihir]

$$CERT[X : Y] = (pk[Y], Y, \dots, \mathcal{S}_{sk[X]}(pk[Y], Y, \dots))$$

To verify *CERT_{Mihir}* you need only *pk_{CA[USA]}*.

Mihir Bellare

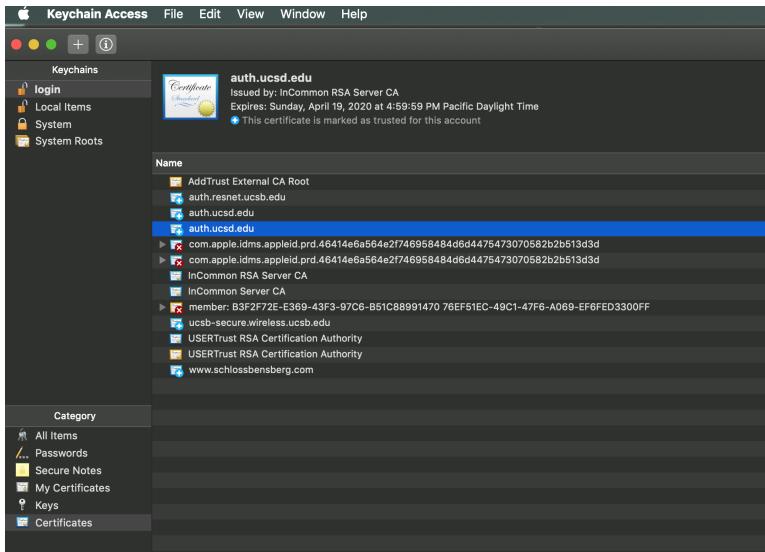
UCSD

21

Why certificate hierarchies?

- It is easier for CA(UCSD) to check Mihir's identity (and issue a certificate) than for CA(USA) since Mihir is on UCSD's payroll and UCSD already has a lot of information about him.
- Spreads the identity-check and certification job to reduce work for individual CAs
- Browsers need to have fewer embedded public keys. (Only root CA public keys needed.)

Certificates on Mac: keychain

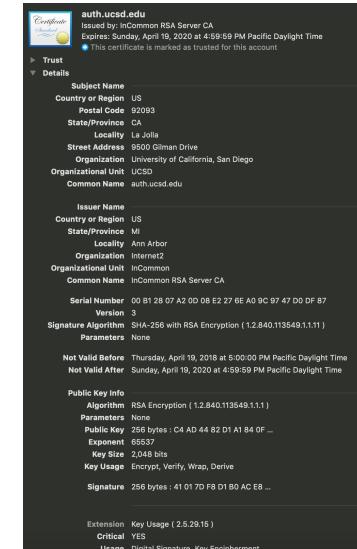


Mihir Bellare

UCSD

23

A particular certificate



Mihir Bellare

UCSD

24

Revocation

Suppose Bob wishes to revoke his certificate $CERT_{Bob}$, perhaps because his secret key was compromised.

- Bob sends $CERT_{Bob}$ and revocation request to CA
- CA checks that request comes from Bob
- CA puts $(CERT_{Bob}, \text{Revocation date})$ on its Certificate Revocation List (CRL)
- This list is disseminated.

Before Alice accepts Bob's certificate she should check that it is not on the CRL.

The OCSP (Online Certificate Status Protocol) is one-way to do this.

Revocation Issues

- November 22: Bob's secret key compromised
- November 24: Bob's $CERT_{Bob}$ revoked
- November 25: Alice sees CRL

In the period Nov. 22-25, $CERT_{Bob}$ might be used and Alice might be accepting as authentic signatures that are really the adversary's. Also Alice might be encrypting data for Bob which the adversary can decrypt.

In practice, CRLs are large and revocation is a problem.

Certificate transparency



Certificate Transparency

Search this site

Documentation

What is Certificate Transparency?

How Certificate Transparency Works

How Log Proofs Work

Benefits and Advantages

Comparison with Other Technologies

Getting Started

Extended Validation in Chrome

Known Logs

General Transparency

Developer Resources

Open Source Project

Certificate Transparency Forum

Certificate Transparency hack days

Certificate Transparency in Chrome

Certificate Transparency in OpenSSL

Resources for site owners

Mailing Lists

Home

Google's Certificate Transparency project fixes several structural flaws in the SSL certificate system, which is the main cryptographic system that underlies all HTTPS connections. These flaws weaken the reliability and effectiveness of encrypted Internet connections and can compromise critical TLS/SSL mechanisms, including domain validation, end-to-end encryption, and the chains of trust set up by certificate authorities. If left unchecked, these flaws can facilitate a wide range of security attacks, such as website spoofing, server impersonation, and man-in-the-middle attacks.



Certificate Transparency helps eliminate these flaws by providing an open framework for monitoring and auditing SSL certificates in nearly real time. Specifically, Certificate Transparency makes it possible to detect SSL certificates that have been mistakenly issued by a certificate authority or maliciously acquired from an otherwise unimpeachable certificate authority. It also makes it possible to identify certificate authorities that have gone rogue and are maliciously issuing certificates.

Because it is an open and public framework, anyone can build or access the basic components that drive Certificate Transparency. This is particularly beneficial to Internet security stakeholders, such as domain owners, certificate authorities, and browser manufacturers, who have a vested interest in maintaining the health and integrity of the SSL certificate system.

Dissemination of public key via server

▲ Not Secure | keyserver.ubuntu.com

SKS OpenPGP Key server

Extract a key

You can find a key by typing in some words that appear in the userid (name, email, etc.) of the key you're looking for, or by typing in the keyid in hex format ("0x...")

Search for a public key

String

Show PGP Fingerprints

Show SKS full-key hashes

Get regular index of matching keys

Get verbose index of matching keys

Retrieve ascii-armored keys

Retrieve keys by full-key hash

Reset Search for a key

Difflie

Search on SGS key-server

Search results for 'difie'

```
Type bits/keyID      cr. time    exp time   key expire
_____
pub 3072R/B8760337 2016-05-24
uid Joe Diffie <bigmondoog1@ yahoo.com>
sig sig3 B8760337 2016-05-24 _____ [selfsig]
sub 3072R/89B6676B 2016-05-24
sig sbind B8760337 2016-05-24 _____ l1
sub 3072R/C1A6676B 2016-05-24
sig sbind B8760337 2016-05-24 _____ l1

pub 2048D/186FEE55A 2015-06-20
uid Manuel Rachad Haj-Saleh Ramirez (Clave Diffie-Hellman) <manuel@hajsaleh.es>
sig sig3 186FEE55A 2015-06-20 _____ 2019-06-20 [selfsig]
sub 2048g/C4F8E190 2015-06-20
sig sbind 186FEE55A 2015-06-20 _____ 2019-06-20 l1

pub 1024D/B1CF6414 2014-12-20
uid diffie-hell-test2048 <diffie-hell-test2048@test.cn>
sig sig3 B1CF6414 2014-12-20 _____ [selfsig]
sub 2048g/0709418D 2014-12-20
sig sbind B1CF6414 2014-12-20 _____ l1

pub 2048R/4BFP2F1E 2012-08-21
uid diffie <diffie@mailivault.com>
sig sig3 4BFP2F1E 2012-08-21 _____ [selfsig]
sub 2048R/5BE9521D 2012-08-21
sig sbind 4BFP2F1E 2012-08-21 _____ l1
```

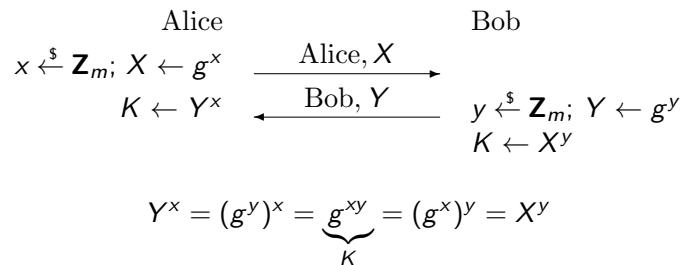
Mihir Bellare

UCSD

29

Diffie-Hellman Key Exchange

Let $G = \langle g \rangle$ be a cyclic group of order m and assume G, g, m are public quantities.



This enables Alice and Bob to agree on a common key K which can subsequently be used, say to encrypt:

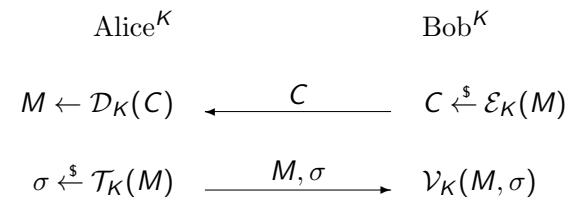
$$\begin{array}{ccc} \text{Alice} & & \text{Bob} \\ M \leftarrow \mathcal{D}_K(C) & \xleftarrow{C} & C \xleftarrow{\$} \mathcal{E}_K(M) \end{array}$$

Mihir Bellare

UCSD

31

Shared key setting



Alice and Bob can

- send each other encrypted data
- verify each other's MACs

Can be preferable to public key setting because computation costs are lower.

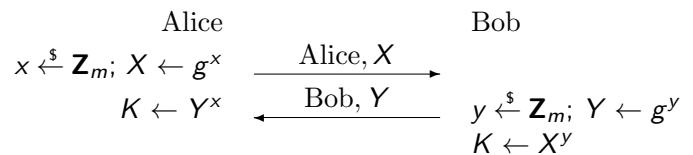
But how do Alice and Bob get a shared key?

Mihir Bellare

UCSD

30

Security of DH Key Exchange under Passive Attack



Eavesdropping adversary gets $X = g^x$ and $Y = g^y$ and wants to compute g^{xy} . But this is the (presumed hard) CDH problem.

Conclusion: DH key exchange is secure against passive (eavesdropping) attack.

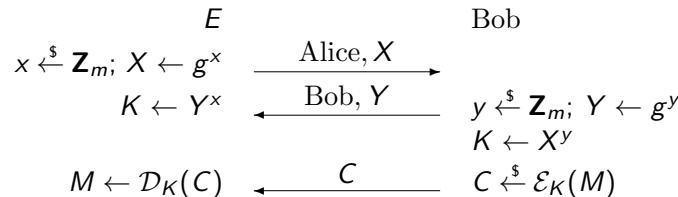
Mihir Bellare

UCSD

32

Security of DH Key Exchange under Active Attack

Entity-in-the-middle attack:



Adversary E impersonates Alice so that:

- Bob thinks he shares K with Alice but E has K
- E can now decrypt ciphertexts Bob intends for Alice

Conclusion: DH key exchange is insecure against active attack

When is key agreement possible?

In the presence of an active adversary, it is impossible for Alice and Bob to

- start from scratch, and
- exchange messages to get a common key unknown to the adversary

Why? Because there is no way for Bob to distinguish Alice from the adversary.

Alice and Bob need some a priori “information advantage” over the adversary. This typically takes the form of long-lived keys.

Settings and long-lived keys

- Public key setting: A has pk_B and B has pk_A
- Symmetric setting: A, B share a key K
- Three party setting: A, B each share a key with a trusted server S .

These keys constitute the long-lived information.

Session keys: The “real” key distribution problem

In practice, Alice and Bob will engage in multiple communication “sessions.” For each, they

- First use a session-key distribution protocol, based on their long-lived keys, to get a fresh, authentic session key;
- Then encrypt or authenticate data under this session key for the duration of the session

Why session keys?

- In public-key setting, efficient cryptography compared to direct use of long-lived keys
- Security attributes, in particular enabling different applications to use keys in different ways and not compromise security of other applications

Session key distribution

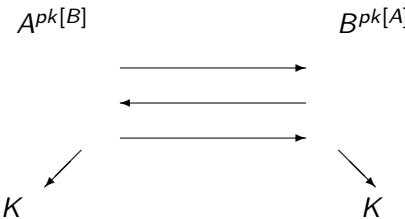
- Hundreds of protocols
- Dozens of security requirements
- Lots of broken protocols
- Protocols easy to specify and hard to get right
- Used ubiquitously: SSL, TLS, SSH, ...

Mihir Bellare

UCSD

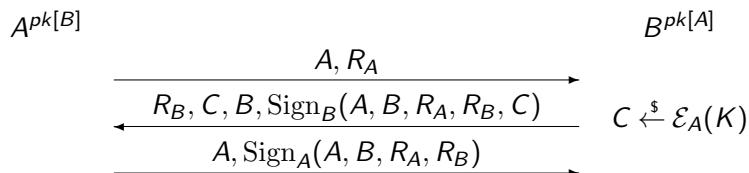
37

Session key exchange in public key setting



Most important type of session key exchange in practice, used in all communication security protocols: SSL, SSH, TLS, IPSEC, 802.11, ...

Protocol KE1: Basic Exchange



Session key K is chosen by B .

R_A, R_B are random nonces.

$\text{Sign}_P(M)$ is P 's signature of M under $sk[P]$. It is verifiable given $pk[P]$.

$\mathcal{E}_A(\cdot)$ is encryption under A 's public key $pk[A]$, decryptable using $sk[A]$.

Mihir Bellare

UCSD

39

Beyond KE2

Modern session key exchange protocols enhance KE2 to incorporate many other security features including

- **Forward security:** Exposure of long-lived secret keys should not compromise PRIOR session keys
- **Anonymity:** Eavesdropper cannot figure out identity (public key) of client authenticating to some server.

Mihir Bellare

UCSD

40

Passwords

A password is a human-memorable key.

Attackers are capable of forming a set D of possible passwords called a dictionary such that

- If the target password pw is in D and
- The attacker knows $\overline{pw} = f(pw)$, the image of pw under some public function f .

then the target password can be found via

for all $pw' \in D$ do
if $f(pw') = \overline{pw}$ then return pw'

This is called a dictionary attack.

Mihir Bellare

UCSD

41

Preventing dictionary attacks: Password selection

Systems try to force users to select “good” passwords, meaning ones not in the dictionary. But studies show that a significant fraction of user passwords end up being in the dictionary anyway.

Attackers get better and better at building dictionaries.

Good password selection helps, but it is unrealistic to think that even the bulk of passwords are well selected, meaning not in the dictionary.

Mihir Bellare

UCSD

43

Password usage

Fact is that in spite of all the great crypto around, a significant fraction of our security today resides in passwords: bank ATM passwords; login passwords; passwords for different websites; ...

Few of us today have cryptographic keys; but we all have more passwords than we can remember!

Passwords are convenient and entrenched.

Preventing dictionary attacks is an important concern.

Mihir Bellare

UCSD

42

Popular passwords

In 2016, the 25 most common passwords made up more than 10% of surveyed passwords, with the most common making up 4%.

Top 25 most common passwords by year according to SplashData							
Rank	2011 ^[4]	2012 ^[5]	2013 ^[6]	2014 ^[7]	2015 ^[8]	2016 ^[9]	2017 ^[9]
1	password	password	123456	123456	123456	123456	123456
2	123456	123456	password	password	password	password	password
3	12345678	12345678	12345678	12345678	12345678	12345678	12345678
4	qwert	abc123	qwert	12345678	qwert	12345678	qwert
5	abc123	qwert	abc123	qwert	12345	football	12345
6	monkey	monkey	123456789	123456789	123456789	qwert	123456789
7	1234567	letmein	111111	1234	football	1234567890	letmein
8	letmein	dragon	1234567	baseball	1234	1234567	sunshine
9	trustno1	iloveyou	dragon	1234567	princess	football	qwert
10	dragon	baseball	adobe123 ^[3]	football	baseball	1234	iloveyou
11	baseball	iloveyou	123123	1234567	welcome	login	admin
12	111111	trustno1	admin	monkey	1234567890	welcome	admin
13	iloveyou	1234567	1234567890	letmein	abc123	solo	monkey
14	master	sunshine	letmein	111111	abc123	login	666666
15	sunshine	master	photoshop ^[3]	111111	1qaz2wsx	admin	abc123
16	ashley	123123	1234	mustang	dragon	121212	starwars
17	bailey	welcome	monkey	access	master	flower	123123
18	passw0rd	shadow	shadow	shadow	monkey	passw0rd	dragon
19	shadow	ashley	sunshine	master	letmein	dragon	passw0rd
20	123123	football	12345	michael	login	sunshine	master
21	654321	jesus	password1	superman	princess	master	hello
22	superman	michael	princess	696969	qwertuiop	hotte	charlie
23	qazwsx	ninja	azerty	123123	solo	loveme	whatever
24	michael	mustang	trustno1	batman	passw0rd	zaq1zaq1	qazwsx
25	Football	password1	000000	trustno1	starwars	password1	trustno1

Mihir Bellare

UCSD

44

Preventing dictionary attacks: avoiding image revelation

An alternative approach is to ensure that usage of a password pw never reveals an image $\bar{pw} = f(pw)$ of pw under a public function f . Then, even if the password is in the dictionary, the dictionary attack cannot be mounted.

Password-based session-key exchange:

- A, B share a password pw .
- They want to interact to get a common session key.
- The protocol should resist dictionary attack: adversary should be unable to obtain an image of pw under a public function.